

⑩ NameNode:

Summary

Security is off.

Safemode is off.

12 files and directories, 5 blocks (5 replicated blocks, 0 erasure coded block groups) = 17 total filesystem object(s).

Heap Memory used 121.06 MB of 431.5 MB Heap Memory. Max Heap Memory is 3.43 GB.

Non Heap Memory used 48.05 MB of 49.19 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	567.68 GB
Configured Remote Capacity:	0 B
DFS Used:	56.23 KB (0%)
Non DFS Used:	356.45 GB
DFS Remaining:	182.25 GB (32.1%)
Block Pool Used:	56.23 KB (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	2 (Decommissioned: 0, In Maintenance: 0)
Dead Nodes	0 (Decommissioned: 0, In Maintenance: 0)
Decommissioning Nodes	0
Entering Maintenance Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	5
Number of Blocks Pending Deletion (including replicas)	0
Block Deletion Start Time	Thu Dec 15 23:51:52 +0300 2022
Last Checkpoint Time	Thu Dec 15 23:48:04 +0300 2022
Enabled Erasure Coding Policies	RS-6-3-1024k

Browse Directory

Go!

Show25 ▾ entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	root	supergroup	0 B	Dec 15 23:52	0	0 B	rmstate	

Showing 1 to 1 of 1 entries

Previous


1

Next

Datanode Information

✓ In service ⚠ Down ⏸ Decommissioning ⚡ Decommissioned ⚰ Decommissioned & dead
🔧 Entering Maintenance ⚙ In Maintenance 🛑 In Maintenance & dead

Datanode usage histogram



In operation

Show 25 entries

Search:

Node	Http Address	Last contact	Last Block Report	Capacity	Blocks	Block pool used	Version
✓ 2d2c07890142:9866 (172.18.0.7:9866)	http://2d2c07890142:9864	2s	5m	283.84 GB	5	28.12 KB (0%)	3.2.1
✓ 8df9ba8504f3:9866 (172.18.0.3:9866)	http://8df9ba8504f3:9864	2s	5m	283.84 GB	5	28.12 KB (0%)	3.2.1

Showing 1 to 2 of 2 entries

[Previous](#) [1](#) [Next](#)

NameNode Journal Status

Current transaction ID: 39

Journal Manager	State
FileJournalManager(root=/hadoop/dfs/name)	EditLogFileOutputStream(/hadoop/dfs/name/current/edits_inprogress_0000000000000000003)

NameNode Storage

Storage Directory	Type	State
/hadoop/dfs/name	IMAGE_AND_EDITS	Active

DFS Storage Types

Storage Type	Configured Capacity	Capacity Used	Capacity Remaining	Block Pool Used	Nodes In Service
DISK	567.68 GB	56.23 KB (0%)	182.24 GB (32.1%)	56.23 KB	2

10 Resource Manager:



All Applic

Google Translate

Logged in as: root

Cluster

About
Nodes
Node Labels
Applications

NEW
NEW SAVING
SUBMITTED
ACCEPTED
RUNNING
FINISHED
FAILED
KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
0	0	0	0	0	0 B	32 GB	0 B	0	16	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
2	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 entries Search:

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCores	Allocated Memory MB	Reserved CPU VCores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
----	------	------	------------------	-------	----------------------	-----------	------------	------------	-------	-------------	--------------------	----------------------	---------------------	---------------------	--------------------	------------	--------------	----------	-------------	-------------------

No data available in table

Showing 0 to 0 of 0 entries

First Previous Next Last



Nodes of the cluster

Logged in as: root

Cluster

About
Nodes
Node Labels
Applications

NEW
NEW SAVING
SUBMITTED
ACCEPTED
RUNNING
FINISHED
FAILED
KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
0	0	0	0	0	0 B	32 GB	0 B	0	16	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
2	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=Mi), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Show 20 entries Search:

Node Labels	Rack	Node State	Node Address	Node HTTP Address	Last health-update	Health-report	Containers	Allocation Tags	Mem Used	Mem Avail	VCores Used	VCores Avail	Version
/default-rack		RUNNING	9585b4648967:43479	9585b4648967:8042	Thu Dec 15 21:28:15 +0000 2022		0		0 B	16 GB	0	8	3.2.1
/default-rack		RUNNING	9eaf0120ad6d:43337	9eaf0120ad6d:8042	Thu Dec 15 21:28:15 +0000 2022		0		0 B	16 GB	0	8	3.2.1

Showing 1 to 2 of 2 entries

First Previous 1 Next Last



NEW,NEW_SAVING,SUBMITTED,ACCEPTED,RUNNING Applications

Logged in as: root

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW_SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved
0	0	0	0	0	0 B	32 GB	0 B	0	16	0

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes	Rebooted Nodes	Shutdown Nodes
2	0	0	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation	Maximum Cluster Application Priority
Capacity Scheduler	[memory-mb (unit=M), vcores]	<memory:1024, vCores:1>	<memory:8192, vCores:4>	0

Dump scheduler logs 1 min

Application Queues

Legend: Capacity Used Used (over capacity) Max Capacity Users Requesting Resources Auto Created Queues

Queue: root 0.0% used

Queue: default 0.0% used

Show 20 entries

Search:

ID	User	Name	Application Type	Queue	Application Priority	StartTime	LaunchTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU VCores	Allocated Memory MB	Reserved CPU VCores	Reserved Memory MB	% of Queue	% of Cluster	Progress	Tracking UI	Blacklisted Nodes
No data available in table																				

Showing 0 to 0 of 0 entries

First Previous Next Last

Aggregate scheduler counts

Total Container Allocations(count)	Total Container Releases(count)	Total Fulfilled Reservations(count)	Total Container Preemptions(count)
0	0	0	0

Last scheduler run

Time	Allocations(count - resources)	Reservations(count - resources)	Releases(count - resources)
Thu Dec 15 21:28:55 +0000 2022	0 - <memory:0, vCores:0>	0 - <memory:0, vCores:0>	0 - <memory:0, vCores:0>

Last Preemption

Time	Container Id	Node Id	Queue
N/A	N/A	N/A	N/A

Last Reservation

Time	Container Id	Node Id	Queue
N/A	N/A	N/A	N/A

Last Allocation

Time	Container Id	Node Id	Queue
N/A	N/A	N/A	N/A

Last Release

Time	Container Id	Node Id	Queue
N/A	N/A	N/A	N/A

FAILED
KILLED

Scheduler

Tools

Dump scheduler logs 1 min

Application Queues

Legend: Capacity Used Used (over capacity) Max Capacity Users Requesting Resources Auto Created Queues

Queue: root 0.0% used
Queue: default 0.0% used

'default' Queue Status

Queue State: RUNNING
Used Capacity: <memory:0, vCores:0> (0.0%)
Configured Capacity: <memory:0, vCores:0>
Configured Max Capacity: unlimited
Effective Capacity: <memory:32768, vCores:16> (100.0%)
Effective Max Capacity: <memory:32768, vCores:16> (100.0%)
Absolute Used Capacity: 0.0%
Absolute Configured Capacity: 100.0%
Absolute Configured Max Capacity: 100.0%
Used Resources: <memory:0, vCores:0>
Configured Max Application Master Limit: 10.0
Max Application Master Resources: <memory:4096, vCores:1>
Used Application Master Resources: <memory:0, vCores:0>
Max Application Master Resources Per User: <memory:4096, vCores:1>
Num Schedulable Applications: 0
Num Non-Schedulable Applications: 0
Num Containers: 0
Max Applications: 10000
Max Applications Per User: 10000
Configured Minimum User Limit Percent: 100%
Configured User Limit Factor: 1.0
Accessible Node Labels: *
Ordering Policy: FifoOrderingPolicy
Preemption: disabled
Intra-queue Preemption: disabled
Default Node Label Expression: <DEFAULT_PARTITION>
Default Application Priority: 0

Active Users Info

User Name	Max Resource	Weight	Used Resource	Max AM Resource	Used AM Resource	Schedulable Apps	Non-Schedulable Apps
-----------	--------------	--------	---------------	-----------------	------------------	------------------	----------------------

⑩ MapReduce for mean value:

```
root@f0b48522740c:/# hadoop jar /opt/hadoop-3.2.1/share/hadoop/tools/lib/hadoop-streaming-3.2.1.jar -file mapper_mean.py -mapper mapper_mean.py -file reducer_mean.py -reducer reducer_mean.py -input /AB_NYC_2019.csv -output mean3
2022-12-18 12:25:17,549 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [mapper_mean.py, reducer_mean.py, /tmp/hadoop-unjar6273384495079626774/] [] /tmp/streamjob421315842509172170.jar tmpDir=null
2022-12-18 12:25:18,178 INFO client.RMProxy: Connecting to ResourceManager at resourcemanager/172.18.0.6:8032
2022-12-18 12:25:18,289 INFO client.AHSProxy: Connecting to Application History server at historyserver/172.18.0.5:10200
2022-12-18 12:25:18,312 INFO client.RMProxy: Connecting to ResourceManager at resourcemanager/172.18.0.6:8032
2022-12-18 12:25:18,312 INFO client.AHSProxy: Connecting to Application History server at historyserver/172.18.0.5:10200
2022-12-18 12:25:18,450 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/root/.staging/job_1671364070864_0005
2022-12-18 12:25:18,529 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:25:18,597 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:25:18,620 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:25:18,673 INFO mapred.FileInputFormat: Total input files to process : 1
2022-12-18 12:25:18,697 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:25:18,713 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:25:18,721 INFO mapreduce.JobSubmitter: number of splits:2
2022-12-18 12:25:18,807 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:25:18,821 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1671364070864_0005
2022-12-18 12:25:18,821 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-12-18 12:25:18,939 INFO conf.Configuration: resource-types.xml not found
2022-12-18 12:25:18,939 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-12-18 12:25:19,197 INFO impl.YarnClientImpl: Submitted application application_1671364070864_0005
2022-12-18 12:25:19,230 INFO mapreduce.Job: The url to track the job: http://resourcemanager:8088/proxy/application_1671364070864_0005/
2022-12-18 12:25:19,231 INFO mapreduce.Job: Running job: job_1671364070864_0005
2022-12-18 12:25:23,300 INFO mapreduce.Job: Job job_1671364070864_0005 running in uber mode : false
2022-12-18 12:25:23,301 INFO mapreduce.Job: map 0% reduce 0%
2022-12-18 12:25:28,366 INFO mapreduce.Job: map 100% reduce 0%
2022-12-18 12:25:32,390 INFO mapreduce.Job: map 100% reduce 100%
2022-12-18 12:25:33,410 INFO mapreduce.Job: Job job_1671364070864_0005 completed successfully
2022-12-18 12:25:33,465 INFO mapreduce.Job: Counters: 54
```

```
root@f0b48522740c:/# hdfs dfs -cat /user/root/mean3/part*
2022-12-18 12:27:57,709 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
152.72068718682894
root@f0b48522740c:/#
```

⑩ MapReduce for var value:

```
root@f0b48522740c:/# hadoop jar /opt/hadoop-3.2.1/share/hadoop/tools/lib/hadoop-streaming-3.2.1.jar -file mapper_var.py -mapper mapper_var.py -file reducer_var.py -reducer reducer_var.py -input /AB_NYC_2019.csv -output var
2022-12-18 12:29:40,781 WARN streaming.StreamJob: -file option is deprecated, please use generic option -files instead.
packageJobJar: [mapper_var.py, reducer_var.py, /tmp/hadoop-unjar5101932596493368928/] [] /tmp/streamjob169141015672428457.jar tmpDir=null
2022-12-18 12:29:41,485 INFO client.RMProxy: Connecting to ResourceManager at resourcemanager/172.18.0.6:8032
2022-12-18 12:29:41,576 INFO client.AHSProxy: Connecting to Application History server at historyserver/172.18.0.5:10200
2022-12-18 12:29:41,601 INFO client.RMProxy: Connecting to ResourceManager at resourcemanager/172.18.0.6:8032
2022-12-18 12:29:41,601 INFO client.AHSProxy: Connecting to Application History server at historyserver/172.18.0.5:10200
2022-12-18 12:29:41,730 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/root/.staging/job_1671364070864_0006
2022-12-18 12:29:41,806 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:29:41,878 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:29:41,894 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:29:41,943 INFO mapred.FileInputFormat: Total input files to process : 1
2022-12-18 12:29:41,963 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:29:41,981 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:29:41,989 INFO mapreduce.JobSubmitter: number of splits:2
2022-12-18 12:29:42,074 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
2022-12-18 12:29:42,086 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1671364070864_0006
2022-12-18 12:29:42,086 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-12-18 12:29:42,206 INFO conf.Configuration: resource-types.xml not found
2022-12-18 12:29:42,207 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-12-18 12:29:42,458 INFO impl.YarnClientImpl: Submitted application application_1671364070864_0006
2022-12-18 12:29:42,487 INFO mapreduce.Job: The url to track the job: http://resourcemanager:8088/proxy/application_1671364070864_0006/
2022-12-18 12:29:42,488 INFO mapreduce.Job: Running job: job_1671364070864_0006
2022-12-18 12:29:46,548 INFO mapreduce.Job: Job job_1671364070864_0006 running in uber mode : false
2022-12-18 12:29:46,550 INFO mapreduce.Job:  map 0% reduce 0%
2022-12-18 12:29:51,610 INFO mapreduce.Job:  map 100% reduce 0%
2022-12-18 12:29:55,635 INFO mapreduce.Job:  map 100% reduce 100%
2022-12-18 12:29:55,644 INFO mapreduce.Job: Job job_1671364070864_0006 completed successfully
2022-12-18 12:29:55,702 INFO mapreduce.Job: Counters: 54
File System Counters
```

```
root@f0b48522740c:/# hdfs dfs -ls /user/root/var
Found 2 items
-rw-r--r-- 3 root supergroup 0 2022-12-18 12:29 /user/root/var/_SUCCESS
-rw-r--r-- 3 root supergroup 19 2022-12-18 12:29 /user/root/var/part-00000
root@f0b48522740c:/# hdfs dfs -cat /user/root/var/part*
2022-12-18 12:32:06,356 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
57675.20659426322
root@f0b48522740c:/#
```